

Retail Inventory Database (CS204 Assignment 3)

--Part 1 #1 Problem:

--Write an SQL statement that calculates total sales

--by product category using an aggregate function.

-- Logic:

--join orders products and categories to tie each other to a product

--and its respective category and then sum order count per category.
The aggregate function is SUM();

```
SELECT
    c.category_name,
    SUM(o.order_amount) AS total_sales
FROM Orders o
    INNER JOIN Products p ON o.product_id = p.product_id
    INNER JOIN Categories c ON p.category_id = c.category_id
GROUP BY c.category_name --we do group by category name so that the sum
is shown per category
ORDER BY total_sales DESC; -- sort from ascending to descending
```

--Result p1#1:

	category_name	total_sales
1	Shoes	90
2	Shirts	60
3	Accessories	50
4	Pants	45

--Part 1 #2

--Write an SQL statement that formats product names
--and supplier contact information using string functions.

--Logic:

-- we join products to suppliers by supplier id with UPPER for product names,
-- CONCAT to build the labels and LOWER for email,
-- which results in one row per product with a description and a supplier contact

```
SELECT
    CONCAT(UPPER(p.product_name), ' - supplied by ', s.supplier_name) AS
product_and_supplier,
    CONCAT('Phone: ', s.phone, ' | Email: ', LOWER(s.email)) AS
supplier_contact
FROM Products p
    INNER JOIN Suppliers s ON p.supplier_id = s.supplier_id;
```

-- Result p1#2:

	□ product_and_supplier ▾	□ supplier_contact ▾
1	COTTON T-SHIRT - supplied by Metro Wear	Phone: +19295551201 Email: metro@shop.com
2	BLUE JEANS - supplied by City Apparel	Phone: +19295551202 Email: city@shop.com
3	RUNNING SHOES - supplied by Prime Clothing	Phone: +19295551204 Email: prime@shop.com
4	LEATHER BELT - supplied by Basics Co	Phone: +19295551205 Email: basics@shop.com

--Part 1 #3

--Write an SQL statement that lists orders by purchase month

--in descending order using date/time functions.

-- Logic : DATE_FORMAT is used to format a full date to year/month and then form there

-- we count how many orders per month and then sum their amounts ordered from latest month to oldest

```
SELECT  
    DATE_FORMAT(order_date, '%Y-%m') AS purchase_month,  
    COUNT(*) AS orders_count,  
    SUM(order_amount) AS total_sales  
FROM Orders  
GROUP BY purchase_month  
ORDER BY purchase_month DESC;
```

-- Result p1#3:

SELECT
 DATE_FORMAT(order_date, '%Y-%m') AS purchase_month,
 COUNT(*) AS orders_count,
 SUM(order_amount) AS total_sales
FROM Orders
GROUP BY purchase_month
ORDER BY purchase_month DESC;

	purchase_month	orders_count	total_sales
1	2025-11	4	245
2	2025-10	2	270
3	2025-09	1	80

--Part 1 #4

--Write an SQL query to calculate the
--20% discounted price for the most expensive product

--Logic: We sort by unitprice from Highest to Lowest and then we only show the first result

-- via LIMIT 1 and retrieve the highest product and then we apply a multiplier to the price

-- to add a 20% discount

```
SELECT
    product_id,
    product_name,
    unit_price,
    unit_price * 0.80 AS discounted_price
FROM Products
ORDER BY unit_price DESC
LIMIT 1;
```

-- Result p1#4:

```
SELECT
    product_id,
    product_name,
    unit_price,
    unit_price * 0.80 AS discounted_price
FROM Products
ORDER BY unit_price DESC
LIMIT 1;
```

The screenshot shows a database interface with two tabs: 'Output' and 'Result 11 ×'. The 'Result' tab displays a table with four columns: product_id, product_name, unit_price, and discounted_price. There is one row of data: product_id 1, product_name 'Running Shoes', unit_price 90, and discounted_price 72.00.

product_id	product_name	unit_price	discounted_price
1	Running Shoes	90	72.00

--Part 2 #1

--Create a view that returns the top 5 best-selling products.

-- Logic : we created a view with a LEFT JOIN of orders so that we can capture all products even the ones with no sales

-- we leverage SUM() for the quantity and order_amount per product to record total quantity sold & total_revenue

-- we limit the view to 5 to show the top 5 products, then we use the SELECT * FROM Top5Products to see the results of the view

```
CREATE VIEW Top5Products AS
SELECT
    p.product_id,
    p.product_name,
    COALESCE(SUM(o.quantity), 0) AS total_quantity_sold,
    COALESCE(SUM(o.order_amount), 0) AS total_revenue
FROM Products p
    LEFT JOIN Orders o ON o.product_id = p.product_id
GROUP BY p.product_id, p.product_name
ORDER BY total_quantity_sold DESC
LIMIT 5;

SELECT * FROM Top5Products;
```

-- Result p2#1:

```
CREATE VIEW Top5Products AS
SELECT
    p.product_id,
    p.product_name,
    COALESCE(SUM(o.quantity), 0) AS total_quantity_sold,
    COALESCE(SUM(o.order_amount), 0) AS total_revenue
FROM Products p
    LEFT JOIN Orders o [l->>0:n] ON o.product_id = p.product_id
GROUP BY p.product_id, p.product_name
ORDER BY total_quantity_sold DESC
LIMIT 5;

SELECT * FROM Top5Products;
```

The screenshot shows a database interface with two tabs: 'Output' and 'Result 15'. The 'Result 15' tab is active and displays a table with four columns: product_id, product_name, total_quantity_sold, and total_revenue. The data is as follows:

	product_id	product_name	total_quantity_sold	total_revenue
1	1	Cotton T-Shirt	7	140
2	2	Blue Jeans	3	135
3	4	Running Shoes	3	270
4	5	Leather Belt	2	50

--Part 2 #2

--Create a stored procedure that accepts a product ID and returns total quantity sold and revenue.

--Then, write a command to execute the stored procedure.

-- Logic : we initialize our procedure with CREATE PROCEDURE and give it a name as GetProductSalesSummary

-- with a required variable int containing the in_product_id , we then use coalesce to safeguard our check for a product with no orders

-- as that scenario would return null, so we return 0 instead. , then besides that logic we just

-- SUM() product order quantity and their respective order_amounts to prepare total_quantity_sold and total_revenue

-- we use delimiter to clearly show the start and end of our statement. and finally we call our procedure with CALL GetProductSalesSummary(1);

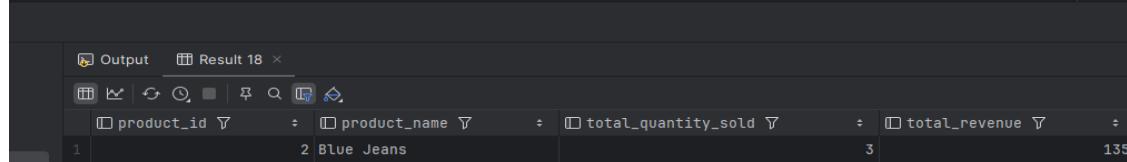
```
DELIMITER $$

CREATE PROCEDURE GetProductSalesSummary(IN in_product_id INT)
BEGIN
    SELECT
        p.product_id,
        p.product_name,
        COALESCE(SUM(o.quantity), 0) AS total_quantity_sold,
        COALESCE(SUM(o.order_amount), 0) AS total_revenue
    FROM Products p
        LEFT JOIN Orders o ON o.product_id = p.product_id
    WHERE p.product_id = in_product_id
    GROUP BY p.product_id, p.product_name;
END$$
DELIMITER ;

CALL GetProductSalesSummary(1);
```

-- Result p2#2:

```
CALL GetProductSalesSummary( in_product_id 2);
```



The screenshot shows the MySQL Workbench interface with the 'Result' tab selected. The results of the stored procedure execution are displayed in a table with four columns: product_id, product_name, total_quantity_sold, and total_revenue. The data row is as follows:

product_id	product_name	total_quantity_sold	total_revenue
1	2 Blue Jeans	3	135

--Part 2 #3

--Part 2 #3 : Create a trigger that logs an entry into an Inventory Audit table every time a product's QOH is updated. Then, add a rule to prevent negative values in the QOH field.

-- Logic: we first create an Inventory_Audit table to log every QOH as that did not exist we then capture the previous QOH and the updated QOH and add a timestamp when the trigger occurs we have a check for if QOH ever becomes negative as that would be an edge case we want covered and our happy path is if the quantity on hand is positive we resume as BAU. results were tested by updating the QOH for a product and inspecting the inventory audit table

```
CREATE TABLE IF NOT EXISTS Inventory_Audit (
    audit_id INT AUTO_INCREMENT PRIMARY KEY,
    product_id INT NOT NULL,
    old_qoh INT,
    new_qoh INT,
    changed_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    FOREIGN KEY (product_id) REFERENCES Products(product_id)
);

DROP TRIGGER IF EXISTS trg_products_qoh_audit;
DELIMITER $$

CREATE TRIGGER trg_products_qoh_audit
    BEFORE UPDATE ON Products
    FOR EACH ROW
BEGIN
    IF NEW.qoh < 0 THEN
        SIGNAL SQLSTATE '45000' SET MESSAGE_TEXT = 'QOH cannot be
negative';
    END IF;

    INSERT INTO Inventory_Audit (product_id, old_qoh, new_qoh)
    VALUES (OLD.product_id, OLD.qoh, NEW.qoh);
END$$
DELIMITER ;

ALTER TABLE Products
    ADD CONSTRAINT chk_products_qoh_nonnegative CHECK (qoh >= 0);
```

	audit_id	product_id	old_qoh	new_qoh	changed_at
1	1	1	1	145	150 2025-12-01 00:26:41

--Part 2 #4

--Create a transaction block that updates inventory and inserts a sales record, rolling back if any part fails.

Logic : we create a function called ProcessSale, in the event of an error we initiate a roll back, we create a new unique order every time ProcessSale is called, we can improve this method by accepting variables so that we do not hardcode it to the same quantity, price, and product each time but this should solve the problem at hand. if there is no error we commit the transaction and do not perform a rollback

```
DROP PROCEDURE IF EXISTS ProcessSale;

DELIMITER $$

CREATE PROCEDURE ProcessSale()
BEGIN

    DECLARE new_order_id INT;

    DECLARE EXIT HANDLER FOR SQLEXCEPTION
        BEGIN
            ROLLBACK;
        END;

    SELECT COALESCE(MAX(order_id), 0) + 1
    INTO new_order_id
    FROM Orders;

    START TRANSACTION;

    UPDATE Products
    SET qoh = qoh - 3
    WHERE product_id = 1;

    INSERT INTO Orders (
        order_id,
        customer_id,
        product_id,
        quantity,
        order_amount,
        order_date,
        order_time
    )
    VALUES (
        new_order_id,
```

-- Result p2#4: showing results for before and after process sale:

```
2
3 ✓ SELECT 'before' AS stage, QOH
4   FROM Products
5 WHERE product_id = 1;
6
7 ✓ CALL ProcessSale();
8
9 ✓ SELECT 'after' AS stage, QOH
10  FROM Products
11 WHERE product_id = 1;
12

Output Result 81 × Result 81-3
□ stage □ QOH
1 before 129
```

```
✓ | SELECT 'after' AS stage, QOH  
| FROM Products  
| WHERE product_id = 1;
```

Output		Result 81	Result 81-3
			X
1	after		126

--Part 2 #5

--Create an index on a frequently queried column and explain why you chose it.

-- Logic : product_id is used in joins and lookups between Orders and Products, so indexing it should help optimize those queries.

```
CREATE INDEX index_orders_product_id ON Orders(product_id);
```